

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>).

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/> (<https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

```
In [1]: from google.colab import drive  
drive.mount('/content/drive')
```

```
o to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee649  
hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Awg%3Aoauth%3A2.0%3Aoob&scope=email%20https%3A%2F%2Fww  
.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive%20https%3A%2F%2Fwww.googleapis.  
com%2Fauth%2Fdrive.photos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fpeopleapi.readonly&response_type=code
```

```
nter your authorization code:  
.....  
ounted at /content/drive
```

[1]. Reading Data

[1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [0]: matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from sklearn.model_selection import GridSearchCV,train_test_split

from sklearn.linear_model import LogisticRegression
```

```
In [0]: using SQLite Table to read data.  
on = sqlite3.connect('database.sqlite')  
  
filtering only positive and negative reviews i.e.  
not taking into consideration those reviews with Score=3  
SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points  
you can change the number to any other number based on your computing power  
  
filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000""", con)  
for tsne assignment you can take 5k data points  
  
filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 5000""", con)  
  
Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0).  
ef partition(x):  
    if x < 3:  
        return 0  
    return 1  
  
changing reviews with score less than 3 to be positive and vice-versa  
ctualScore = filtered_data['Score']  
ositiveNegative = actualScore.map(partition)  
filtered_data['Score'] = positiveNegative  
rint("Number of data points in our data", filtered_data.shape)  
filtered_data.head(3)
```

Number of data points in our data (5000, 10)

Out[0]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality Dog Food
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertised
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all



```
In [0]: isplay = pd.read_sql_query("""
    SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
    FROM Reviews
    GROUP BY UserId
    HAVING COUNT(*)>1
    """", con)
```

```
In [0]: print(display.shape)
display.head()

(80668, 7)
```

Out[0]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

```
In [0]: display[display['UserId']=='AZY10LLTJ71NX']
```

Out[0]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...	5

```
In [0]: display['COUNT(*)'].sum()
```

Out[0]: 393063

[2] Exploratory Data Analysis

[2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [0]: isplay= pd.read_sql_query("""
ELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURN"
ORDER BY ProductID
""", con)
isplay.head()
```

Out[0]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Sun
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADF VANILL WAFER
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADF VANILL WAFER
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADF VANILL WAFER
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADF VANILL WAFER
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACK QUADF VANILL WAFER



As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [0]: Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_position='last')
```

```
In [0]: Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=False)
final.shape
```

```
Out[0]: (4986, 10)
```

```
In [0]: Checking to see how much % of data still remains
final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[0]: 99.72
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [0]: isplay= pd.read_sql_query("""
ELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

isplay.head()
```

Out[0]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1	5	1224892800	Bought This for My Son College
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2	4	1212883200	Pure cocoa taste wi crunchy almond inside

```
In [0]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [0]: Before starting the next phase of preprocessing lets see the number of entries left  
rint(final.shape)  
  
How many positive and negative reviews are present in our dataset?  
inal['Score'].value_counts()  
  
(4986, 10)  
  
Out[0]: 1    4178  
0     808  
Name: Score, dtype: int64
```

[3] Preprocessing

[3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [0]: printing some random reviews
ent_0 = final['Text'].values[0]
print(sent_0)
print("=-*50)

ent_1000 = final['Text'].values[1000]
print(sent_1000)
print("=-*50)

ent_1500 = final['Text'].values[1500]
print(sent_1500)
print("=-*50)

ent_4900 = final['Text'].values[4900]
print(sent_4900)
print("=-*50)
```

hy is this \$[...] when the same product is available for \$[...] here?
<http://www.amazon.com/VICTOR-FLY-MAGNET-BIT-REFILL/dp/B00004RBDY>

The Victor M380 and M502 traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

=====

recently tried this flavor/brand and was surprised at how delicious these chips are. The best thing was that there were a lot of "brown" chips in the bsg (my favorite), so I bought some more through amazon and shared with family and friends. I am a little disappointed that there are not, so far, very many brown chips in these bags, but the flavor is still very good. I like them better than the yogurt and green onion flavor because they do not seem to be as salted, and the onion flavor is better. If you haven't eaten Kettle chips before, I recommend that you try a bag before buying bulk. They are thicker and crunchier than Lays but just as fresh out of the bag.

=====

ow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other wants crispy cookie. Hey, I'm sorry; but these reviews do nobody any good beyond reminding us to look before ordering.

The ones are chocolate-oatmeal cookies. If you don't like that combination, don't order this type of cookie. I find the combo quite nice, really. The oatmeal sort of "calms" the rich chocolate flavor and gives the cookie sort of a coconut-type consistency. Now let's also remember that tastes differ; so, I've given my opinion.

Then, these are soft, chewy cookies -- as advertised. They are not "crispy" cookies, or the blurb would say "crispy," rather than "chewy." I happen to like raw cookie dough; however, I don't see where these taste like raw cookie dough. Both are soft, however, so is this the confusion? And, yes, they stick together. Soft cookies tend to do that. They aren't individually wrapped, which would add to the cost. Oh yeah, chocolate chip cookies tend to be somewhat sweet.

So, if you want something hard and crisp, I suggest Nabisco's Ginger Snaps. If you want a cookie that's soft, chewy and tastes like a combination of chocolate and oatmeal, give these a try. I'm here to place my second order.

=====

ove to order my coffee on amazon. easy and shows up quickly.
This k cup is great coffee. dcaf is very good as well

```
In [0]: remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_150)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

hy is this \$[...] when the same product is available for \$[...] here?

The Victor M380 and M502 traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

In [0]: <https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-tags-from-an-element>

```
from bs4 import BeautifulSoup

oup = BeautifulSoup(sent_0, 'lxml')
ext = soup.get_text()
rint(text)
rint("=*50)

oup = BeautifulSoup(sent_1000, 'lxml')
ext = soup.get_text()
rint(text)
rint("=*50)

oup = BeautifulSoup(sent_1500, 'lxml')
ext = soup.get_text()
rint(text)
rint("=*50)

oup = BeautifulSoup(sent_4900, 'lxml')
ext = soup.get_text()
rint(text)
```

hy is this \$[...] when the same product is available for \$[...] here? />The Victor M380 and M502 traps are unreal, o
course -- total fly genocide. Pretty stinky, but only right nearby.

=====

recently tried this flavor/brand and was surprised at how delicious these chips are. The best thing was that there
ere a lot of "brown" chips in the bsg (my favorite), so I bought some more through amazon and shared with family and
riends. I am a little disappointed that there are not, so far, very many brown chips in these bags, but the flavor
s still very good. I like them better than the yogurt and green onion flavor because they do not seem to be as salt
, and the onion flavor is better. If you haven't eaten Kettle chips before, I recommend that you try a bag before b
ying bulk. They are thicker and crunchier than Lays but just as fresh out of the bag.

=====

ow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other wants crispy cookie
. Hey, I'm sorry; but these reviews do nobody any good beyond reminding us to look before ordering.These are choco
ate-oatmeal cookies. If you don't like that combination, don't order this type of cookie. I find the combo quite n
ice, really. The oatmeal sort of "calms" the rich chocolate flavor and gives the cookie sort of a coconut-type consi
istency. Now let's also remember that tastes differ; so, I've given my opinion.Then, these are soft, chewy cookies --
s advertised. They are not "crispy" cookies, or the blurb would say "crispy," rather than "chewy." I happen to lik
raw cookie dough; however, I don't see where these taste like raw cookie dough. Both are soft, however, so is this
he confusion? And, yes, they stick together. Soft cookies tend to do that. They aren't individually wrapped, whic
would add to the cost. Oh yeah, chocolate chip cookies tend to be somewhat sweet.So, if you want something hard an
crisp, I suggest Nabisco's Ginger Snaps. If you want a cookie that's soft, chewy and tastes like a combination of c
ocolate and oatmeal, give these a try. I'm here to place my second order.

=====

ove to order my coffee on amazon. easy and shows up quickly.This k cup is great coffee. dcaf is very good as well

```
In [0]: https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"\n\t", " not", phrase)
    phrase = re.sub(r"\re", " are", phrase)
    phrase = re.sub(r"\s", " is", phrase)
    phrase = re.sub(r"\d", " would", phrase)
    phrase = re.sub(r"\ll", " will", phrase)
    phrase = re.sub(r"\t", " not", phrase)
    phrase = re.sub(r"\ve", " have", phrase)
    phrase = re.sub(r"\m", " am", phrase)
    return phrase
```

```
In [0]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("=="*50)
```

ow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other wants crispy cookie . Hey, I am sorry; but these reviews do nobody any good beyond reminding us to look before ordering.

Th se are chocolate-oatmeal cookies. If you do not like that combination, do not order this type of cookie. I find th combo quite nice, really. The oatmeal sort of "calms" the rich chocolate flavor and gives the cookie sort of a coc nut-type consistency. Now let is also remember that tastes differ; so, I have given my opinion.

Then, th se are soft, chewy cookies -- as advertised. They are not "crispy" cookies, or the blurb would say "crispy," rather han "chewy." I happen to like raw cookie dough; however, I do not see where these taste like raw cookie dough. Bot are soft, however, so is this the confusion? And, yes, they stick together. Soft cookies tend to do that. They a e not individually wrapped, which would add to the cost. Oh yeah, chocolate chip cookies tend to be somewhat sweet. br />
So, if you want something hard and crisp, I suggest Nabisco is Ginger Snaps. If you want a cookie that is oft, chewy and tastes like a combination of chocolate and oatmeal, give these a try. I am here to place my second o der.

=====

```
In [0]: remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

hy is this \$ [...] when the same product is available for \$ [...] here?
 />
The Victor and traps are unrea , of course -- total fly genocide. Pretty stinky, but only right nearby.

```
In [0]: remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

ow So far two two star reviews One obviously had no idea what they were ordering the other wants crispy cookies Hey am sorry but these reviews do nobody any good beyond reminding us to look before ordering br br These are chocolate atmeal cookies If you do not like that combination do not order this type of cookie I find the combo quite nice real y The oatmeal sort of calms the rich chocolate flavor and gives the cookie sort of a coconut type consistency Now le is also remember that tastes differ so I have given my opinion br br Then these are soft chewy cookies as advertise They are not crispy cookies or the blurb would say crispy rather than chewy I happen to like raw cookie dough howev r I do not see where these taste like raw cookie dough Both are soft however so is this the confusion And yes they s ick together Soft cookies tend to do that They are not individually wrapped which would add to the cost Oh yeah choc late chip cookies tend to be somewhat sweet br br So if you want something hard and crisp I suggest Nabisco is Ginger naps If you want a cookie that is soft chewy and tastes like a combination of chocolate and oatmeal give these a try am here to place my second order

In [0]:

<https://gist.github.com/sebleier/554280>

we are removing the words from the stop words list: 'no', 'nor', 'not'

 => after the above steps, we are getting "br br"
we are including them into stop words list
instead of
 if we have
 these tags would have removed in the 1st step

```
topwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", \
    "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', \
    'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their', \
    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those', \
    'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does', \
    'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of', \
    'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', 'before', 'after', \
    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again', 'furthe
    ', \
    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'mor
    ', \
    'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very', \
    's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're',
    \
    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn',
    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't", 'mustn',
    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'weren',
    "wer
    't", \
    'won', "won't", 'wouldn', "wouldn't"])
```

In [0]: *Combining all the above students*

```
from tqdm import tqdm
reprocessed_reviews = []
tqdm is for printing the status bar
for sentance in tqdm(final['Text'].values):
    sentance = re.sub(r"http\S+", "", sentance)
    sentance = BeautifulSoup(sentance, 'lxml').get_text()
    sentance = decontracted(sentance)
    sentance = re.sub("\S*\d\S*", "", sentance).strip()
    sentance = re.sub('[^A-Za-z]+', ' ', sentance)
    # https://gist.github.com/sebleier/554280
    sentance = ' '.join(e.lower() for e in sentance.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentance.strip())
```

100% |██████████| 4986/4986 [00:01<00:00, 3137.37it/s]

In [0]: `reprocessed_reviews[1500]`

Out[0]: 'wow far two two star reviews one obviously no idea ordering wants crispy cookies hey sorry reviews nobody good beyond reminding us look ordering chocolate oatmeal cookies not like combination not order type cookie find combo quite nice really oatmeal sort calms rich chocolate flavor gives cookie sort coconut type consistency let also remember taste differ given opinion soft chewy cookies advertised not crispy cookies blurb would say crispy rather chewy happen like raw cookie dough however not see taste like raw cookie dough soft however confusion yes stick together soft cookie tend not individually wrapped would add cost oh yeah chocolate chip cookies tend somewhat sweet want something hard crisp suggest nabisco ginger snaps want cookie soft chewy tastes like combination chocolate oatmeal give try place second order'

[3.2] Preprocessing Review Summary

In [0]: *# Similarly you can do preprocessing for review summary also.*

[4] Featurization

```
In [0]: ith open('/content/drive/My Drive/Applied AI assignments/preprocessed_reviews.pkl','rb') as f:  
    preprocessed_reviews=pickle.load(f)  
ith open('/content/drive/My Drive/Applied AI assignments/list_of_scores.pkl','rb') as f:  
    list_of_scores=pickle.load(f)
```

[4.1] BAG OF WORDS

```
In [0]: Bow  
Count_vect = CountVectorizer() #in scikit-learn  
Count_vect.fit(preprocessed_reviews)  
print("some feature names ", count_vect.get_feature_names()[:10])  
print('*'*50)  
  
final_counts = count_vect.transform(preprocessed_reviews)  
print("the type of count vectorizer ",type(final_counts))  
print("the shape of out text BOW vectorizer ",final_counts.get_shape())  
print("the number of unique words ", final_counts.get_shape()[1])  
  
some feature names  ['aa', 'aahhhs', 'aback', 'abandon', 'abates', 'abbott', 'abby', 'abdominal', 'abiding', 'abilit'  
']  
=====  
the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>  
the shape of out text BOW vectorizer  (4986, 12997)  
the number of unique words  12997
```

[4.2] Bi-Grams and n-Grams.

In [0]: bi-gram, tri-gram and n-gram

```
removing stop words like "not" should be avoided before building n-grams
count_vect = CountVectorizer(ngram_range=(1,2))
please do read the CountVectorizer documentation http://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html
```

```
you can choose these numbers min_df=10, max_features=5000, of your choice
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_bigram_counts))
print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_bigram_counts.get_shape()[1])
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144
```

[4.3] TF-IDF

In [0]:

```
f_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
f_idf_vect.fit(preprocessed_reviews)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names()[0:10])
print('='*50)

final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_tf_idf))
print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[1])

some sample features(unique words in the corpus) ['ability', 'able', 'able find', 'able get', 'absolute', 'absolutel',
', 'absolutely delicious', 'absolutely love', 'absolutely no', 'according']
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144
```

[4.4] Word2Vec

```
In [0]: Train your own Word2Vec model using your own text corpus  
=0  
ist_of_sentance=[]  
or sentance in preprocessed_reviews:  
list_of_sentance.append(sentance.split())
```

In [0]: Using Google News Word2Vectors

in this project we are using a pretrained model by google its 3.3G file, once you load this into your memory it occupies ~9Gb, so please do this step only if you have >12G of ram we will provide a pickle file which contains a dict , and it contains all our corpus words as keys and model[word] as values To use this code-snippet, download "GoogleNews-vectors-negative300.bin" from <https://drive.google.com/file/d/0B7XkCwpI5KDYNLNUTLSS21pQmM/edit> it's 1.9GB in size.

*<http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzzPY>
you can comment this whole cell
or change these variable according to your need*

```
s_your_ram_gt_16g=False
ant_to_use_google_w2v = False
ant_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentece,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))

if want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin', binary=True)
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have google's word2vec file, keep want_to_train_w2v = True, to train your own w2v ")
```

```
('snack', 0.9951335191726685), ('calorie', 0.9946465492248535), ('wonderful', 0.9946032166481018), ('excellent', 0.944332838058472), ('especially', 0.9941144585609436), ('baked', 0.9940600395202637), ('salted', 0.994047224521637), ('alternative', 0.9937226176261902), ('tasty', 0.9936816692352295), ('healthy', 0.9936649799346924)]  
===== ('varieties', 0.9994194507598877), ('become', 0.9992934465408325), ('popcorn', 0.9992750883102417), ('de', 0.999261016140747), ('miss', 0.9992451071739197), ('melitta', 0.999218761920929), ('choice', 0.9992102384567261), ('america', 0.9991837739944458), ('beef', 0.9991780519485474), ('finish', 0.9991567134857178)]
```

```
In [0]: w2v_words = list(w2v_model.wv.vocab)  
print("number of words that occurred minimum 5 times ",len(w2v_words))  
print("sample words ", w2v_words[0:50])
```

```
umber of words that occurred minimum 5 times 3817  
ample words  ['product', 'available', 'course', 'total', 'pretty', 'stinky', 'right', 'nearby', 'used', 'ca', 'not', 'beat', 'great', 'received', 'shipment', 'could', 'hardly', 'wait', 'try', 'love', 'call', 'instead', 'removed', 'eas  
ly', 'daughter', 'designed', 'printed', 'use', 'car', 'windows', 'beautifully', 'shop', 'program', 'going', 'lot', 'fun', 'everywhere', 'like', 'tv', 'computer', 'really', 'good', 'idea', 'final', 'outstanding', 'window', 'everybod  
'ask', 'bought', 'made']
```

[4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

In [0]:

```

average Word2Vec
compute average word2vec for each review.
ent_vectors = [] # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentece): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

```

100% |██████████| 4986/4986 [00:03<00:00, 1330.47it/s]

4986

50

[4.4.1.2] TFIDF weighted W2v

In [0]:

```

S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
f_idf_matrix = model.fit_transform(preprocessed_reviews)
# we are converting a dictionary with word as a key, and the idf as a value
idictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

In [0]:

```
TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = [] # the tfidf-w2v for each sentence/review is stored in this list
ow=0;
or sent in tqdm(list_of_senteance): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1
```

100% |██████████| 4986/4986 [00:20<00:00, 245.63it/s]

[5] Assignment 5: Apply Logistic Regression

1. Apply Logistic Regression on these feature sets

- SET 1:Review text, preprocessed one converted into vectors using (BOW)
- SET 2:Review text, preprocessed one converted into vectors using (TFIDF)
- SET 3:Review text, preprocessed one converted into vectors using (AVG W2v)
- SET 4:Review text, preprocessed one converted into vectors using (TFIDF W2v)

2. Hyper parameter tuning (find best hyper parameters corresponding the algorithm that you choose)

- Find the best hyper parameter which will give the maximum AUC (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/receiver-operating-characteristic-curve-roc-curve-and-auc-1/>) value
- Find the best hyper parameter using k-fold cross validation or simple cross validation data
- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter tuning

3. Perturbation Test

- Get the weights W after fit your model with the data X i.e Train data.
- Add a noise to the X ($X' = X + e$) and get the new data set X' (if X is a sparse matrix, $X.data+=e$)
- Fit the model again on data X' and get the weights W'
- Add a small eps value(to eliminate the divisible by zero error) to W and W' i.e $W=W+10^{-6}$ and $W'=W'+10^{-6}$
- Now find the % change between W and W' ($|(W-W')| / |W| * 100$)
- Calculate the 0th, 10th, 20th, 30th, ...100th percentiles, and observe any sudden rise in the values of percentage_change_vector
- Ex: consider your 99th percentile is 1.3 and your 100th percentiles are 34.6, there is sudden rise from 1.3 to 34.6, now calculate the 99.1, 99.2, 99.3,..., 100th percentile values and get the proper value after which there is sudden rise the values, assume it is 2.5
- Print the feature names whose % change is more than a threshold x(in our example it's 2.5)

4. Sparsity

- Calculate sparsity on weight vector obtained after using L1 regularization

NOTE: Do sparsity and multicollinearity for any one of the vectorizers. Bow or tf-idf is recommended.

5. Feature importance

- Get top 10 important features for both positive and negative classes separately.

6. Feature engineering

- To increase the performance of your model, you can also experiment with feature engineering like :
 - Taking length of reviews as another feature.
 - Considering some features from review summary as well.

7. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure.



Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.



Along with plotting ROC curve, you need to print the confusion matrix (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/>) with predicted and original labels of test data points. Please visualize your confusion matrices using seaborn heatmaps.



(<https://seaborn.pydata.org/generated/seaborn.heatmap.html>).

(<https://seaborn.pydata.org/generated/seaborn.heatmap.html>).

(<https://seaborn.pydata.org/generated/seaborn.heatmap.html>).

8. Conclusion (<https://seaborn.pydata.org/generated/seaborn.heatmap.html>)

(<https://seaborn.pydata.org/generated/seaborn.heatmap.html>).

- You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library (<https://seaborn.pydata.org/generated/seaborn.heatmap.html>) link (<http://zetcode.com/python/prettytable/>).



Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on your train data, and apply the method transform() on cv/test data.
4. For more details please go through this link. (<https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf>)

Applying Logistic Regression

[5.1] Logistic Regression on BOW, SET 1

[5.1.1] Applying Logistic Regression with L1 regularization on BOW, SET 1

```
In [0]: =preprocessed_reviews  
=list_of_scores
```

```
In [0]: split the dataset into train and test sets  
_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=0)
```

```
In [0]: vectorizing using Bow  
bowVectorizer=CountVectorizer()  
bowVectorizer.fit(X_train)  
_train_BOW=bowVectorizer.transform(X_train)  
_test_BOW=bowVectorizer.transform(X_test)
```

```
In [0]: Using GridSearchCV for finding the best C  
rModel=LogisticRegression(penalty='l1')  
hyper_parameters={'C':[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]}  
odel=GridSearchCV(lrModel,hyper_parameters,scoring='roc_auc',cv=3)  
odel.fit(X_train_BOW,y_train)  
  
odel.best_estimator_  
odel.best_score_
```

```
Out[0]: 0.9318178140362692
```

```
In [0]: odel.best_estimator_
```

```
Out[0]: LogisticRegression(C=1, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='warn',
                           n_jobs=None, penalty='l1', random_state=None, solver='warn',
                           tol=0.0001, verbose=0, warm_start=False)
```

In [0]: plotting the train AUC vs CV AUC for different $\log(C)$ values.
 \log is considered because, different scale of values can be displayed on the plot
 $C_values=[10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}, 10^0, 10^1, 10^2, 10^3, 10^4]$
 $train_auc= model.cv_results_['mean_train_score']$
 $v_auc = model.cv_results_['mean_test_score']$

```
=np.log10(C_values)

lt.plot(C, train_auc, label='Train AUC')
    this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

lt.plot(C, cv_auc, label='CV AUC')
    this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
lt.legend()
lt.xlabel("log(C)")
lt.ylabel("AUC")
lt.title("ERROR PLOTS")
lt.show()
```

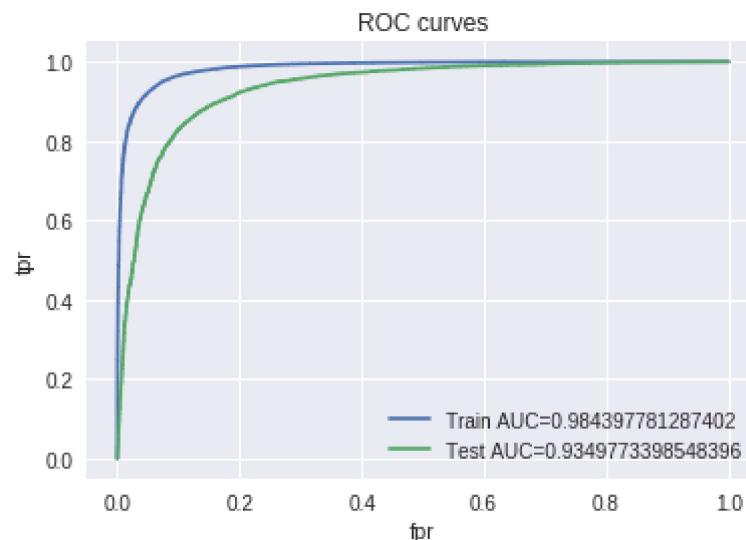


In [0]: est_C=1

```
In [0]: training the model with best C and plotting ROC curves on both train and test data
rModel=LogisticRegression(penalty='l1',C=best_C)
rModel.fit(X_train_BOW,y_train)

train_fpr,train_tpr,thresholds=roc_curve(y_train,lrModel.predict_proba(X_train_BOW)[:,1])
test_fpr,test_tpr,thresholds=roc_curve(y_test,lrModel.predict_proba(X_test_BOW)[:,1])

lt.plot(train_fpr,train_tpr,label='Train AUC='+str(auc(train_fpr,train_tpr)))
lt.plot(test_fpr,test_tpr,label='Test AUC='+str(auc(test_fpr,test_tpr)))
lt.legend()
lt.xlabel('fpr')
lt.ylabel('tpr')
lt.title('ROC curves')
lt.show()
```



In [0]: confusion matrix using seaborn heatmap

```
_pred=lrModel.predict(X_test_BOW)
m=confusion_matrix(y_test,y_pred)
sns.heatmap(cm,annot=True,fmt='.5g')
```

Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd197f66cc0>



[5.1.1.1] Calculating sparsity on weight vector obtained using L1 regularization on BOW, SET 1

In [0]: calculating sparsity on weight vector obtained using L1 regularization on Bow

```
eight_vector=lrModel.coef_
total_no_of_elem=weight_vector.size
nonzero_elem=np.count_nonzero(weight_vector)
sparsity=(total_no_of_elem-nonzero_elem)/total_no_of_elem
print('The sparsity of weight vector obtained using L1 regularization is {}'.format(sparsity))
```

The sparsity of weight vector obtained using L1 regularization is 0.8997330232958705

[5.1.2] Applying Logistic Regression with L2 regularization on BOW, SET 1

```
In [0]: using GridSearchCV to determine the best CV for Logistic regression with L2 regularization
rModel=LogisticRegression(penalty='l2')
hyper_parameters={'C':[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]}
odel=GridSearchCV(lrModel,hyper_parameters,scoring='roc_auc',cv=3)
odel.fit(X_train_BOW,y_train)
```

```
Out[0]: GridSearchCV(cv=3, error_score='raise-deprecating',
estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='warn',
n_jobs=None, penalty='l2', random_state=None, solver='warn',
tol=0.0001, verbose=0, warm_start=False),
fit_params=None, iid='warn', n_jobs=None,
param_grid={'C': [0.0001, 0.001, 0.01, 0.1, 1, 10, 100, 1000, 10000]},
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring='roc_auc', verbose=0)
```

```
In [0]: odel.best_estimator_
```

```
Out[0]: LogisticRegression(C=0.1, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='warn',
n_jobs=None, penalty='l2', random_state=None, solver='warn',
tol=0.0001, verbose=0, warm_start=False)
```

```
In [0]: odel.best_score_
```

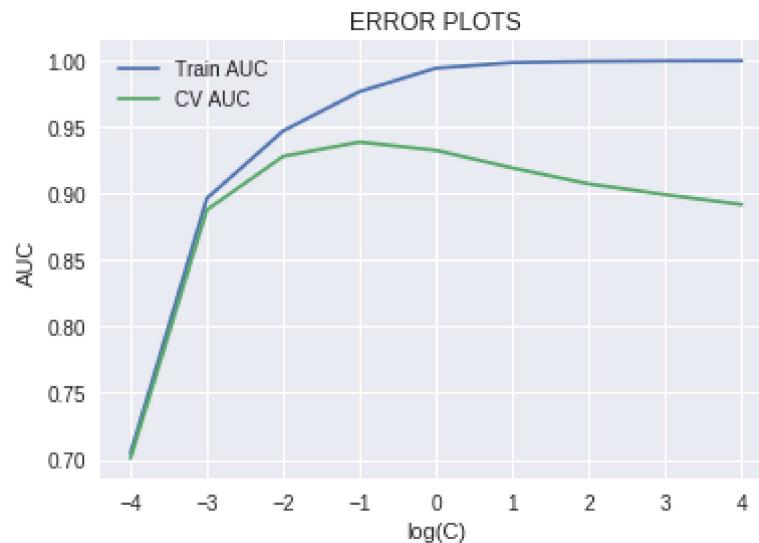
```
Out[0]: 0.9388322723684638
```

```
In [0]: train AUC vs test AUC
_c_values=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
train_auc= model.cv_results_['mean_train_score']
v_auc = model.cv_results_['mean_test_score']

=np.log10(C_values)

lt.plot(C, train_auc, label='Train AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

lt.plot(C, cv_auc, label='CV AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values, cv_auc - cv_auc_std, cv_auc + cv_auc_std, alpha=0.2, color='darkorange')
lt.legend()
lt.xlabel("log(C)")
lt.ylabel("AUC")
lt.title("ERROR PLOTS")
lt.show()
```



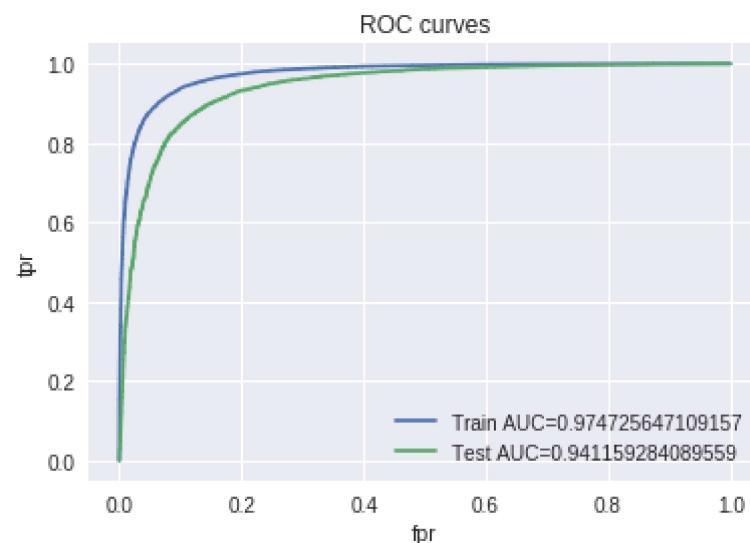
```
In [0]: est_C=0.1
```

In [75]: training the model with best and plotting ROC curves for train and test data

```
rModel=LogisticRegression(penalty='l2',C=best_C)
rModel.fit(X_train_BOW,y_train)

train_fpr,train_tpr,thresholds=roc_curve(y_train,lrModel.predict_proba(X_train_BOW)[:,1])
test_fpr,test_tpr,thresholds=roc_curve(y_test,lrModel.predict_proba(X_test_BOW)[:,1])

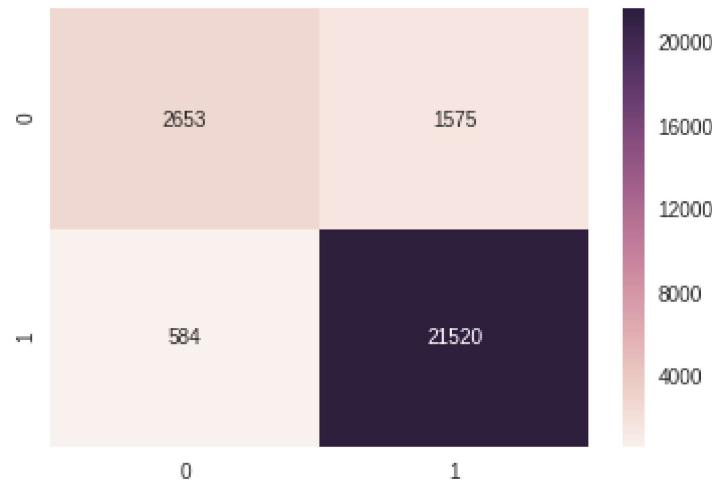
lt.plot(train_fpr,train_tpr,label='Train AUC='+str(auc(train_fpr,train_tpr)))
lt.plot(test_fpr,test_tpr,label='Test AUC='+str(auc(test_fpr,test_tpr)))
lt.legend()
lt.xlabel('fpr')
lt.ylabel('tpr')
lt.title('ROC curves')
lt.show()
```



In [0]: confusion matrix using seaborn heatmap

```
_pred=lrModel.predict(X_test_BOW)
m=confusion_matrix(y_test,y_pred)
ns.heatmap(cm,annot=True,fmt='.5g')
```

Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4568eb18d0>



[5.1.2.1] Performing perturbation test (multicollinearity check) on BOW, SET 1

In [0]: weight vector w
=lrModel.coef_

In [0]: generate noise from normal distribution (mean=0 and std=0.1)
=np.random.normal(loc=0.0, scale=0.1, size=None)

In [0]: adding the error to the train dataset
_train_BOW=X_train_BOW.astype('float64')
_train_BOW.data+=e

```
In [80]: training the model on new dataset (with noise)
rModel=LogisticRegression(penalty='l2',C=best_C)
rModel.fit(X_train_BOW,y_train)
```

```
Out[80]: LogisticRegression(C=0.1, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='warn',
                           n_jobs=None, penalty='l2', random_state=None, solver='warn',
                           tol=0.0001, verbose=0, warm_start=False)
```

```
In [0]: new weight vector w_new
_wnew=lrModel.coef_
```

```
In [0]: adding a small eps value to w and w_new to avoid division by zero error
=w+10**-5
_wnew=w_new+10**-5
```

```
In [0]: finding the percentage change of weight vectors
iff_w_and_wnew=np.subtract(w,w_new)
ercent_change_in_weights=np.absolute(np.divide(iff_w_and_wnew,w))*100
```

```
In [0]: calculating percentiles 0,10,20,30,.....100th percentile
ercentiles=[x for x in range(0,110,10)]
ercentile_values=[]
or p in percentiles:
percentile_values.append(np.percentile(ercent_change_in_weights,p))
```

```
In [91]: ercentile_values
```

```
Out[91]: [0.0001185713175136682,
          0.9846196282428765,
          1.9933017214623678,
          3.0482663111395114,
          4.118409590879273,
          5.301226907020326,
          6.570835871392785,
          8.123048369118022,
          10.31358485948647,
          14.82530743051162,
          89722.9282600134]
```

```
In [0]: ercentile_values_2=[]
or p in range(90,101):
    percentile_values_2.append(np.percentile(percent_change_in_weights,p))
```

```
In [93]: ercentile_values_2
```

```
Out[93]: [14.82530743051162,
          15.62371325662323,
          16.656428526141813,
          17.92805599247506,
          19.480114184036122,
          21.3808656442978,
          24.587104508894043,
          30.726848430526168,
          42.24108169009264,
          73.39763953172897,
          89722.9282600134]
```

```
In [0]: threshold value for obtainiing the features
threshold=73.39
```

```
In [0]: get the feature names
features=bowVectorizer.get_feature_names()
```

```
In [0]: https://docs.scipy.org/doc/numpy/reference/generated/numpy.argwhere.html  
ndices_array=np.argwhere(percent_change_in_weights>=threshold)  
ndices=[]  
or i in range(ndices_array.shape[0]):  
    indices.append(ndices_array[i][1])
```

```
In [0]: features with %change in weights more than threshold  
features=np.array(features)  
features_above_threshold=list(features[indices])
```

```
In [120]: printing only first 10 features with %change in weights more than threshold  
features_above_threshold[:10]
```

```
Out[120]: ['abilities',  
          'acted',  
          'adjusted',  
          'advantage',  
          'affected',  
          'agenda',  
          'aggravated',  
          'aggrevated',  
          'alcoholic',  
          'alkaloid']
```

[5.1.3] Feature Importance on BOW, SET 1

[5.1.3.1] Top 10 important features of positive class from SET 1

```
In [0]: get features
         features=bowVectorizer.get_feature_names()
         get weight vector as list
         weight_vector=lrModel.coef_.tolist()[0]
         get the dictionary with key as feature and its weight as value
         eat_weight_dict=dict(zip(features,weight_vector))
```

```
In [0]: https://www.python.org/dev/peps/pep-0265/
         eat_weight_dict_items=[(v,k) for k,v in feat_weight_dict.items()]
         eat_weight_dict_items.sort()
         eat_weight_dict_items.reverse()
```

```
In [123]: op10_imp_feat_pos_cls=[k for v,k in feat_weight_dict_items][:10]
           op10_imp_feat_pos_cls
```

```
Out[123]: ['delicious',
           'perfect',
           'excellent',
           'wonderful',
           'awesome',
           'great',
           'yummy',
           'amazing',
           'best',
           'hooked']
```

[5.1.3.2] Top 10 important features of negative class from SET 1

```
In [124]: features_ordered_by_weights=[k for v,k in feat_weight_dict_items]
          features_ordered_by_weights.reverse()
          op10_imp_feat_neg_cls=features_ordered_by_weights[:10]
          op10_imp_feat_neg_cls
```

```
Out[124]: ['worst',
           'disappointing',
           'terrible',
           'disappointment',
           'awful',
           'rip',
           'threw',
           'horrible',
           'tasteless',
           'sorry']
```

[5.2] Logistic Regression on TFIDF, SET 2

[5.2.1] Applying Logistic Regression with L1 regularization on TFIDF, SET 2

```
In [0]: Tfidf vectorization
        fidf_vectorizer=TfidfVectorizer(ngram_range=(1,2),min_df=10)
        fidf_vectorizer.fit(X_train)
        _train_tfidf=fidf_vectorizer.transform(X_train)
        _test_tfidf=fidf_vectorizer.transform(X_test)
```

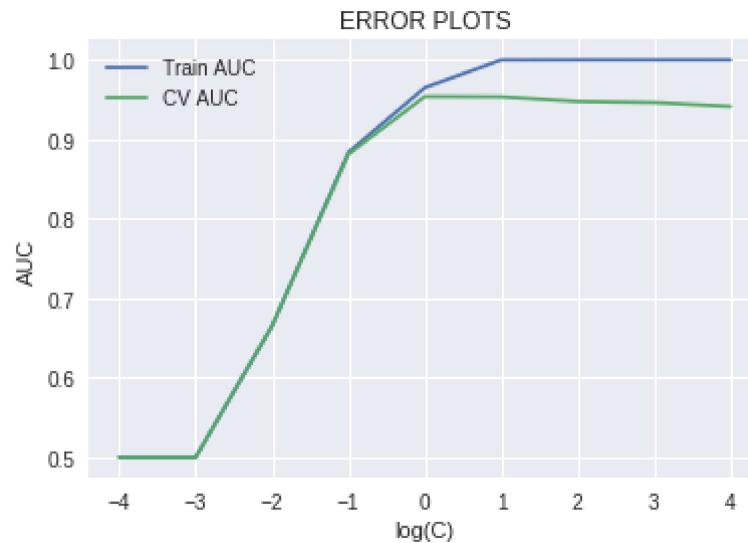
```
In [0]: determining the best C using GridSearchCV
rModel=LogisticRegression(penalty='l1')
hyper_parameters={'C':[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]}
odel=GridSearchCV(lrModel,hyper_parameters,scoring='roc_auc',cv=3)
odel.fit(X_train_tfidf,y_train)

_values=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
train_auc= model.cv_results_['mean_train_score']
v_auc = model.cv_results_['mean_test_score']

=np.log10(C_values)

lt.plot(C, train_auc, label='Train AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

lt.plot(C, cv_auc, label='CV AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values, cv_auc - cv_auc_std, cv_auc + cv_auc_std, alpha=0.2, color='darkorange')
lt.legend()
lt.xlabel("log(C)")
lt.ylabel("AUC")
lt.title("ERROR PLOTS")
lt.show()
```



```
In [0]: odel.best_estimator_
```

```
Out[0]: LogisticRegression(C=1, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='warn',
                           n_jobs=None, penalty='l1', random_state=None, solver='warn',
                           tol=0.0001, verbose=0, warm_start=False)
```

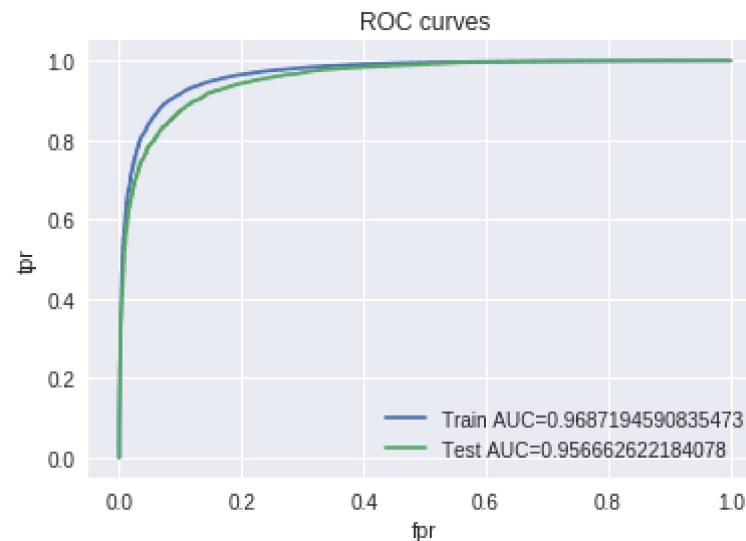
```
In [0]: est_C=1
```

In [0]:

```
Training model on best C
ROC curves for train and test data
rModel=LogisticRegression(penalty='l1',C=best_C)
rModel.fit(X_train_tfidf,y_train)

train_fpr,train_tpr,thresholds=roc_curve(y_train,lrModel.predict_proba(X_train_tfidf)[:,1])
est_fpr,test_tpr,thresholds=roc_curve(y_test,lrModel.predict_proba(X_test_tfidf)[:,1])

lt.plot(train_fpr,train_tpr,label='Train AUC='+str(auc(train_fpr,train_tpr)))
lt.plot(test_fpr,test_tpr,label='Test AUC='+str(auc(test_fpr,test_tpr)))
lt.legend()
lt.xlabel('fpr')
lt.ylabel('tpr')
lt.title('ROC curves')
lt.show()
```



In [0]: confusion matrix using seaborn heatmap

```
_pred=lrModel.predict(X_test_tfidf)
m=confusion_matrix(y_test,y_pred)
ns.heatmap(cm,annot=True,fmt='.5g')
```

Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f29d48a40b8>



[5.2.2] Applying Logistic Regression with L2 regularization on TFIDF, SET 2

In [0]:

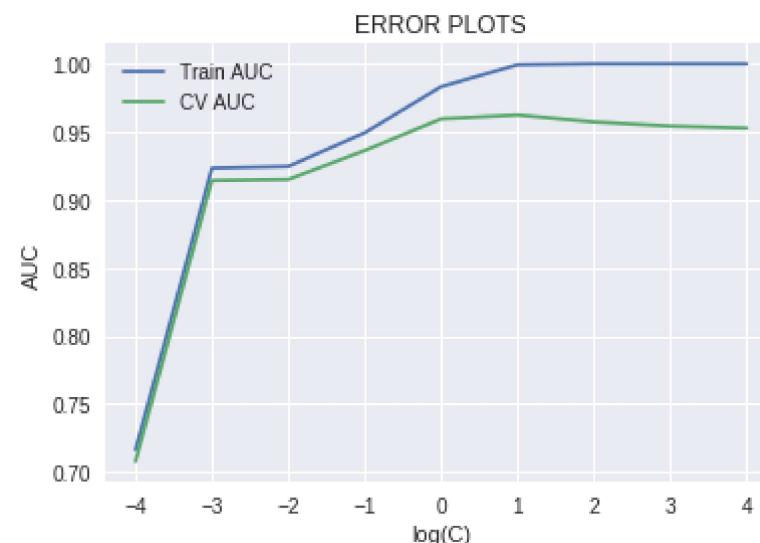
```
lrModel=LogisticRegression(penalty='l2')
hyper_parameters={'C':[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]}
odel=GridSearchCV(lrModel,hyper_parameters,scoring='roc_auc',cv=3)
odel.fit(X_train_tfidf,y_train)

_C_values=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
train_auc= model.cv_results_['mean_train_score']
cv_auc = model.cv_results_['mean_test_score']

=np.log10(C_values)

lt.plot(C, train_auc, label='Train AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

lt.plot(C, cv_auc, label='CV AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
lt.legend()
lt.xlabel("log(C)")
lt.ylabel("AUC")
lt.title("ERROR PLOTS")
lt.show()
```



```
In [0]: odel.best_estimator_
```

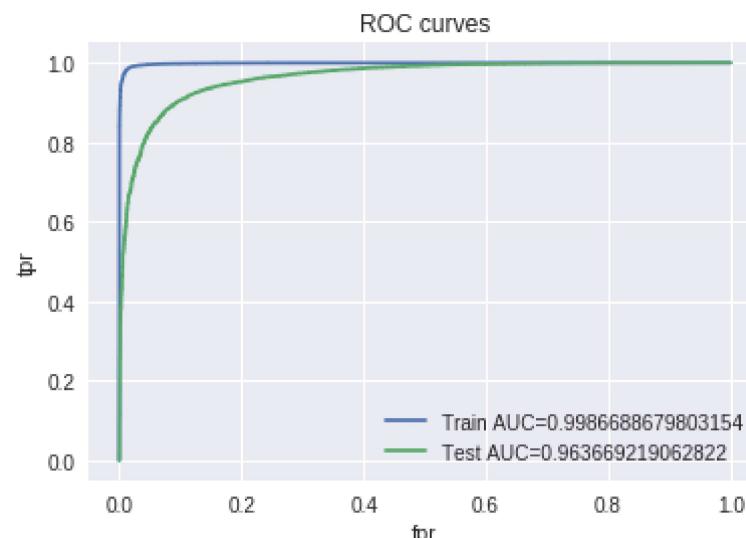
```
Out[0]: LogisticRegression(C=10, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='warn',
                           n_jobs=None, penalty='l2', random_state=None, solver='warn',
                           tol=0.0001, verbose=0, warm_start=False)
```

```
In [0]: est_C=10
```

```
In [0]: rModel=LogisticRegression(penalty='l2',C=best_C)
rModel.fit(X_train_tfidf,y_train)

train_fpr,train_tpr,thresholds=roc_curve(y_train,lrModel.predict_proba(X_train_tfidf)[:,1])
est_fpr,test_tpr,thresholds=roc_curve(y_test,lrModel.predict_proba(X_test_tfidf)[:,1])

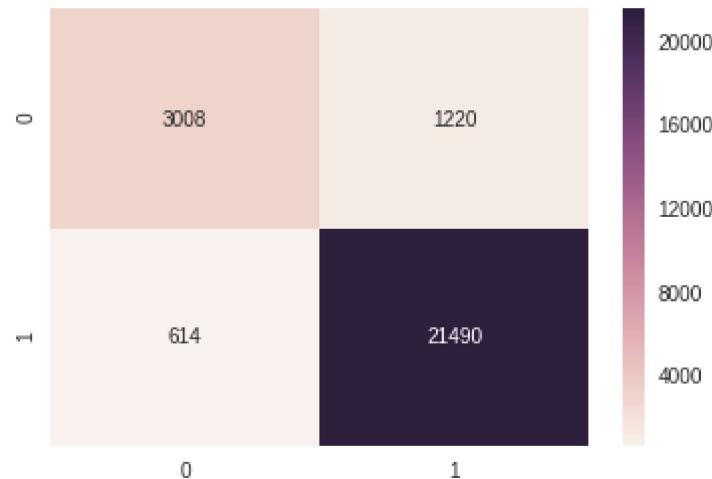
lt.plot(train_fpr,train_tpr,label='Train AUC='+str(auc(train_fpr,train_tpr)))
lt.plot(test_fpr,test_tpr,label='Test AUC='+str(auc(test_fpr,test_tpr)))
lt.legend()
lt.xlabel('fpr')
lt.ylabel('tpr')
lt.title('ROC curves')
lt.show()
```



In [0]: confusion matrix using seaborn heatmap

```
_pred=lrModel.predict(X_test_tfidf)
m=confusion_matrix(y_test,y_pred)
ns.heatmap(cm,annot=True,fmt='.5g')
```

Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f29d8c12400>



[5.2.3] Feature Importance on TFIDF, SET 2

[5.2.3.1] Top 10 important features of positive class from SET 2

In [0]: get features
 features=tfidf_vectorizer.get_feature_names()
 get weight vector as list
 weight_vector=lrModel.coef_.tolist()[0]
 get the dictionary with key as feature and its weight as value
 eat_weight_dict=dict(zip(features,weight_vector))

```
In [0]: https://www.python.org/dev/peps/pep-0265/
eat_weight_dict_items=[(v,k) for k,v in feat_weight_dict.items()]
eat_weight_dict_items.sort()
eat_weight_dict_items.reverse()
```

```
In [0]: op10_imp_feat_pos_cls=[k for v,k in feat_weight_dict_items][:10]
op10_imp_feat_pos_cls
```

```
Out[0]: ['great',
'delicious',
'not disappointed',
'best',
'perfect',
'good',
'excellent',
'nice',
>wonderful',
'loves']
```

[5.2.3.2] Top 10 important features of negative class from SET 2

```
In [0]: eatures_ordered_by_weights=[k for v,k in feat_weight_dict_items]
eatures_ordered_by_weights.reverse()
op10_imp_feat_neg_cls=features_ordered_by_weights[:10]
op10_imp_feat_neg_cls
```

```
Out[0]: ['worst',
'not worth',
'disappointed',
'not recommend',
'disappointing',
'not good',
'terrible',
'disappointment',
'awful',
'weak']
```

[5.3] Logistic Regression on AVG W2V, SET 3

[5.3.1] Applying Logistic Regression with L1 regularization on AVG W2V SET 3

```
In [0]: list of tokenized sentences
```

```
=0
```

```
ist_of_sentance=[]
```

```
or sentance in preprocessed_reviews:
```

```
list_of_sentance.append(sentance.split())
```

```
In [0]: _train,X_test,y_train,y_test=train_test_split(list_of_sentance,y,test_size=0.3,random_state=0)
```

```
In [0]: w2v model on train data
```

```
_train_w2v_model=Word2Vec(X_train,min_count=5,size=50,workers=4)
```

In [0]:

```
vectorizing X_train using Avg W2V
_x_train_w2v_words = list(X_train_w2v_model.wv.vocab)
_x_train_sent_vectors = [] # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(X_train): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in X_train_w2v_words:
            vec = X_train_w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    X_train_sent_vectors.append(sent_vec)

vectorizing X_test using Avg W2V
X_test_w2v_words = list(X_test_w2v_model.wv.vocab)
_x_test_sent_vectors = [] # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(X_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in X_train_w2v_words:
            vec = X_train_w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    X_test_sent_vectors.append(sent_vec)
```

100% |██████████| 61441/61441 [02:13<00:00, 459.92it/s]

100% |██████████| 26332/26332 [00:56<00:00, 463.79it/s]

In [0]:

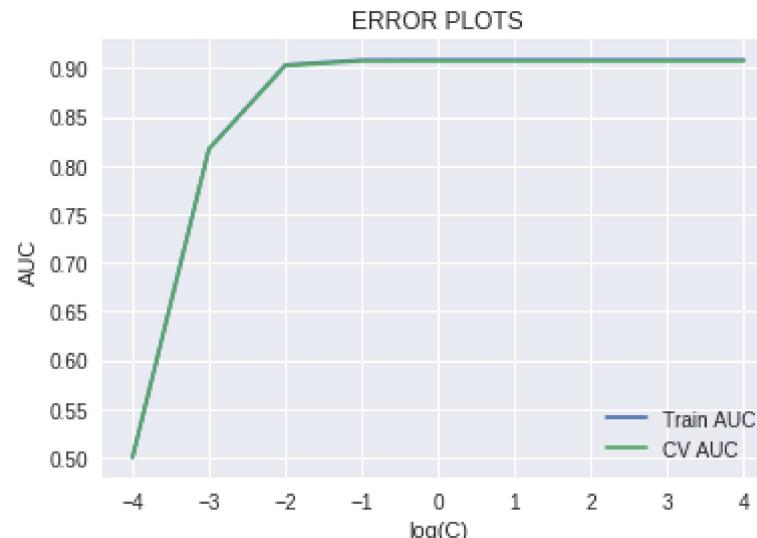
```
rModel=LogisticRegression(penalty='l1')
hyper_parameters={'C':[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]}
odel=GridSearchCV(lrModel,hyper_parameters,scoring='roc_auc',cv=3)
odel.fit(X_train_sent_vectors,y_train)

_C_values=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
train_auc= model.cv_results_['mean_train_score']
v_auc = model.cv_results_['mean_test_score']

=np.log10(C_values)

lt.plot(C, train_auc, label='Train AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

lt.plot(C, cv_auc, label='CV AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
lt.legend()
lt.xlabel("log(C)")
lt.ylabel("AUC")
lt.title("ERROR PLOTS")
lt.show()
```



```
In [0]: odel.best_estimator_
```

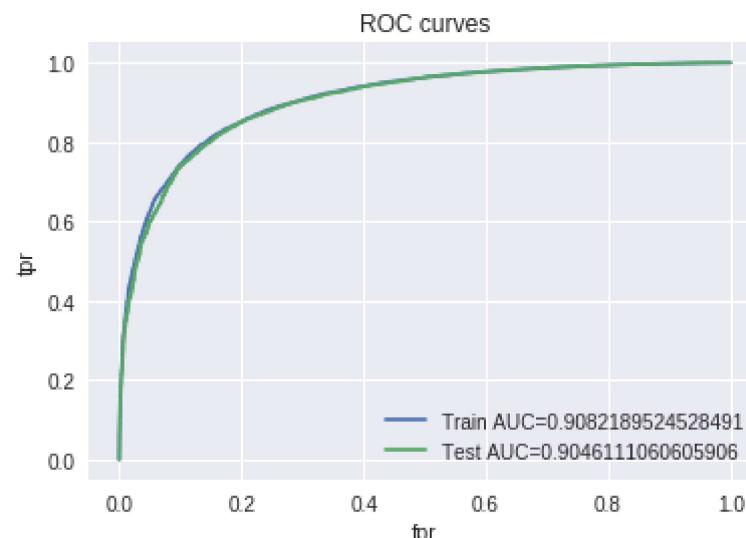
```
Out[0]: LogisticRegression(C=1, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='warn',
                           n_jobs=None, penalty='l1', random_state=None, solver='warn',
                           tol=0.0001, verbose=0, warm_start=False)
```

```
In [0]: est_C=1
```

```
In [0]: rModel=LogisticRegression(penalty='l1',C=best_C)
rModel.fit(X_train_sent_vectors,y_train)

train_fpr,train_tpr,thresholds=roc_curve(y_train,lrModel.predict_proba(X_train_sent_vectors)[:,1])
est_fpr,test_tpr,thresholds=roc_curve(y_test,lrModel.predict_proba(X_test_sent_vectors)[:,1])

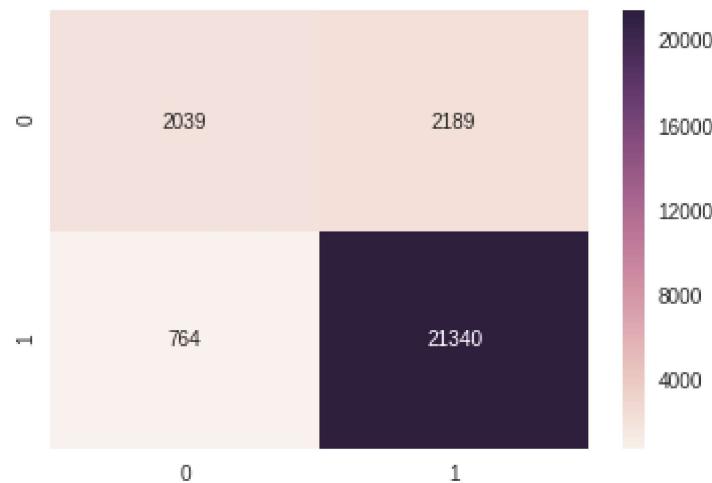
lt.plot(train_fpr,train_tpr,label='Train AUC='+str(auc(train_fpr,train_tpr)))
lt.plot(test_fpr,test_tpr,label='Test AUC='+str(auc(test_fpr,test_tpr)))
lt.legend()
lt.xlabel('fpr')
lt.ylabel('tpr')
lt.title('ROC curves')
lt.show()
```



In [0]: confusion matrix using seaborn heatmap

```
_pred=lrModel.predict(X_test_sent_vectors)
m=confusion_matrix(y_test,y_pred)
sns.heatmap(cm,annot=True,fmt='.5g')
```

Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f29c28c19b0>



[5.3.2] Applying Logistic Regression with L2 regularization on AVG W2V, SET 3

In [0]:

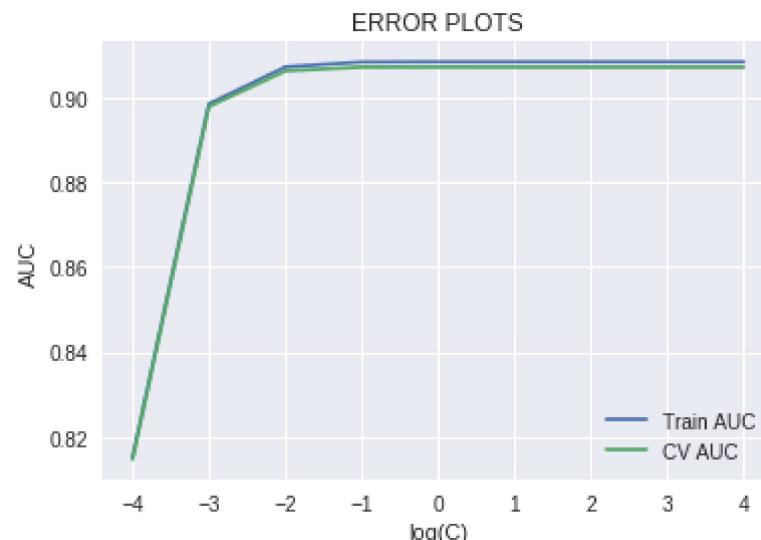
```
rModel=LogisticRegression(penalty='l2')
hyper_parameters={'C':[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]}
odel=GridSearchCV(lrModel,hyper_parameters,scoring='roc_auc',cv=3)
odel.fit(X_train_sent_vectors,y_train)

_C_values=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
train_auc= model.cv_results_['mean_train_score']
v_auc = model.cv_results_['mean_test_score']

=np.log10(C_values)

lt.plot(C, train_auc, label='Train AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

lt.plot(C, cv_auc, label='CV AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
lt.legend()
lt.xlabel("log(C)")
lt.ylabel("AUC")
lt.title("ERROR PLOTS")
lt.show()
```



```
In [0]: odel.best_estimator_
```

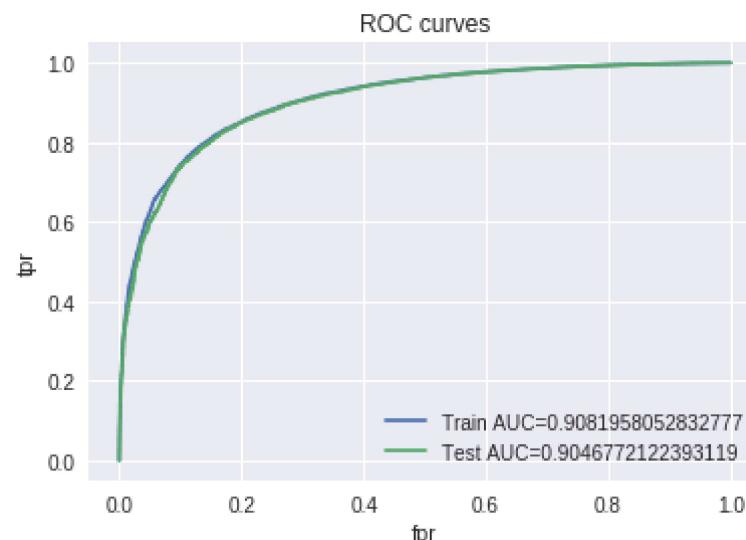
```
Out[0]: LogisticRegression(C=0.1, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='warn',
                           n_jobs=None, penalty='l2', random_state=None, solver='warn',
                           tol=0.0001, verbose=0, warm_start=False)
```

```
In [0]: est_C=0.1
```

```
In [0]: rModel=LogisticRegression(penalty='l2',C=best_C)
rModel.fit(X_train_sent_vectors,y_train)

train_fpr,train_tpr,thresholds=roc_curve(y_train,lrModel.predict_proba(X_train_sent_vectors)[:,1])
est_fpr,test_tpr,thresholds=roc_curve(y_test,lrModel.predict_proba(X_test_sent_vectors)[:,1])

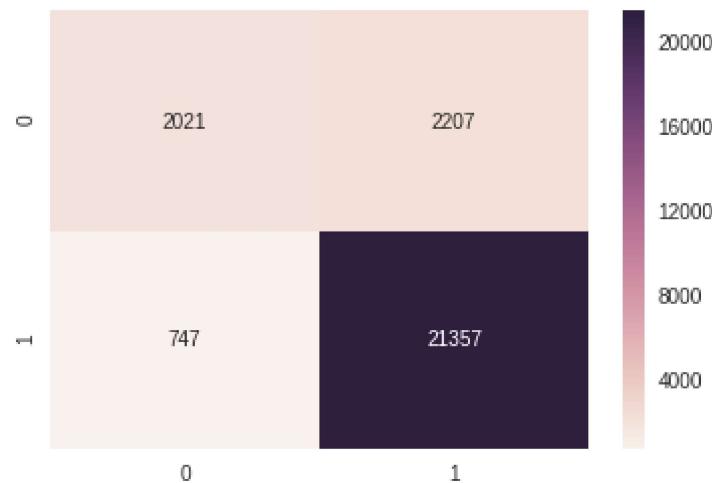
lt.plot(train_fpr,train_tpr,label='Train AUC='+str(auc(train_fpr,train_tpr)))
lt.plot(test_fpr,test_tpr,label='Test AUC='+str(auc(test_fpr,test_tpr)))
lt.legend()
lt.xlabel('fpr')
lt.ylabel('tpr')
lt.title('ROC curves')
lt.show()
```



In [0]: confusion matrix using seaborn heatmap

```
_pred=lrModel.predict(X_test_sent_vectors)
m=confusion_matrix(y_test,y_pred)
ns.heatmap(cm,annot=True,fmt='.5g')
```

Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f29c27a8320>



[5.4] Logistic Regression on TFIDF W2V, SET 4

[5.4.1] Applying Logistic Regression with L1 regularization on TFIDF W2V, SET 4

In [0]: `_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=0)`

In [0]: `odel=TfidfVectorizer()
odel.fit(X_train)
_train_tfidf=odel.transform(X_train)
_test_tfidf=odel.transform(X_test)`

```
In [0]: building a dictionary with features as keys and their corresponding idf's as values
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

```
In [0]: vectorizing X_train using tfidf w2v
_X_train_list_of_sents=[]
or sent in X_train:
_X_train_list_of_sents.append(sent.split())

build w2v model for X_train
_X_train_w2v_model=Word2Vec(X_train_list_of_sents,min_count=5,size=50,workers=4)
ords=list(_X_train_w2v_model.wv.vocab)

TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

_X_train_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
ow=0;
or sent in tqdm(X_train_list_of_sents): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in words and word in tfidf_feat:
            vec = X_train_w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        if weight_sum != 0:
            sent_vec /= weight_sum
    _X_train_tfidf_sent_vectors.append(sent_vec)
    row += 1
```

100% |██████████| 61441/61441 [22:37<00:00, 45.25it/s]

In [0]: vectorizing test dataset

```
_test_list_of_sents=[]
for sent in X_test:
    X_test_list_of_sents.append(sent.split())

build w2v model for X_test
X_test_w2v_model=Word2Vec(X_test_list_of_sents,min_count=5,size=50,workers=4)
X_test_words=list(X_test_w2v_model.wv.vocab)

TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

_test_tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
ow=0;
for sent in tqdm(X_test_list_of_sents): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in words and word in tfidf_feat:
            vec = X_train_w2v_model.wv[word]
            #tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole courpus
            # sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    X_test_tfidf_sent_vectors.append(sent_vec)
    row += 1
```

100% |██████████| 26332/26332 [09:33<00:00, 45.91it/s]

In [0]:

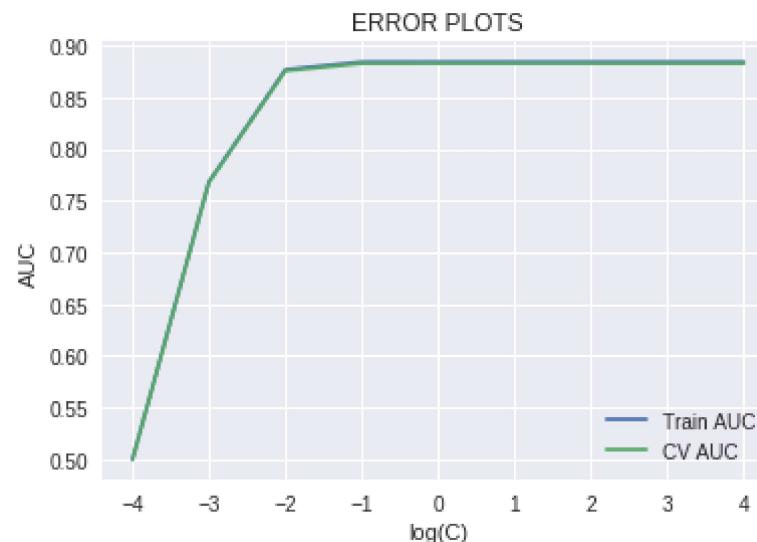
```
lrModel=LogisticRegression(penalty='l1')
hyper_parameters={'C':[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]}
odel=GridSearchCV(lrModel,hyper_parameters,scoring='roc_auc',cv=3)
odel.fit(X_train_tfidf_sent_vectors,y_train)

_C_values=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
train_auc= model.cv_results_['mean_train_score']
cv_auc = model.cv_results_['mean_test_score']

=np.log10(C_values)

lt.plot(C, train_auc, label='Train AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

lt.plot(C, cv_auc, label='CV AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
lt.legend()
lt.xlabel("log(C)")
lt.ylabel("AUC")
lt.title("ERROR PLOTS")
lt.show()
```



```
In [0]: odel.best_estimator_
```

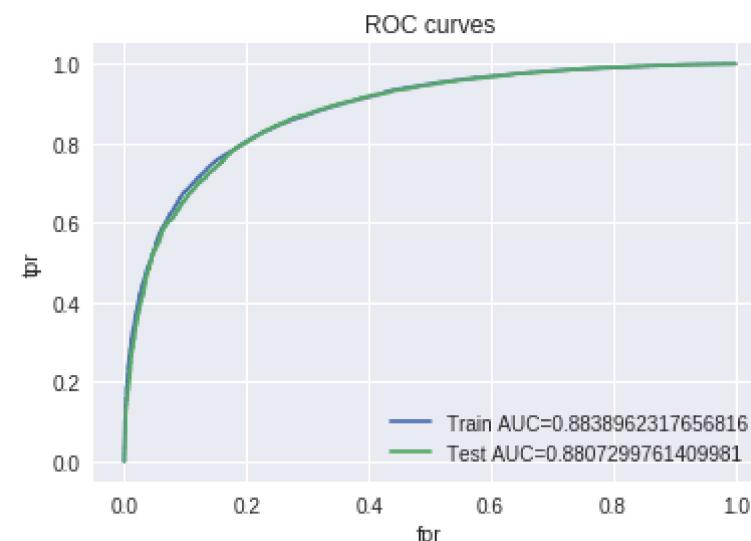
```
Out[0]: LogisticRegression(C=1, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='warn',
                           n_jobs=None, penalty='l1', random_state=None, solver='warn',
                           tol=0.0001, verbose=0, warm_start=False)
```

```
In [0]: est_C=1
```

```
In [0]: rModel=LogisticRegression(penalty='l1',C=best_C)
rModel.fit(X_train_tfidf_sent_vectors,y_train)

train_fpr,train_tpr,thresholds=roc_curve(y_train,lrModel.predict_proba(X_train_tfidf_sent_vectors)[:,1])
est_fpr,test_tpr,thresholds=roc_curve(y_test,lrModel.predict_proba(X_test_tfidf_sent_vectors)[:,1])

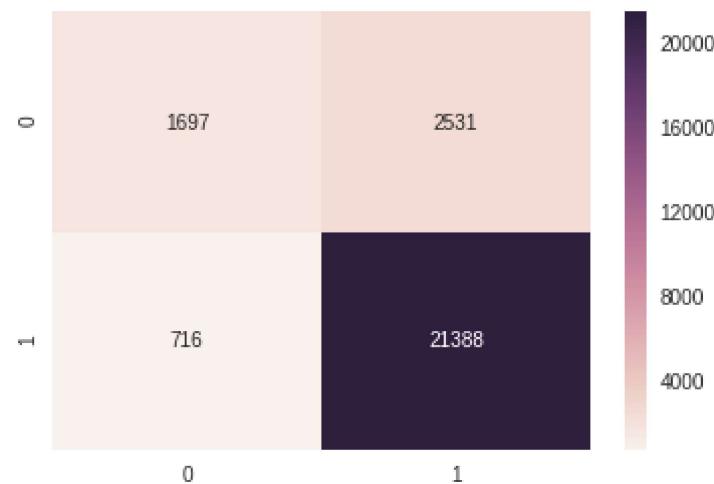
lt.plot(train_fpr,train_tpr,label='Train AUC='+str(auc(train_fpr,train_tpr)))
lt.plot(test_fpr,test_tpr,label='Test AUC='+str(auc(test_fpr,test_tpr)))
lt.legend()
lt.xlabel('fpr')
lt.ylabel('tpr')
lt.title('ROC curves')
lt.show()
```



In [0]: confusion matrix using seaborn heatmap

```
_pred=lrModel.predict(X_test_tfidf_sent_vectors)
m=confusion_matrix(y_test,y_pred)
ns.heatmap(cm,annot=True,fmt='.5g')
```

Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff8e9dcacc0>



[5.4.2] Applying Logistic Regression with L2 regularization on TFIDF W2V, SET 4

In [0]:

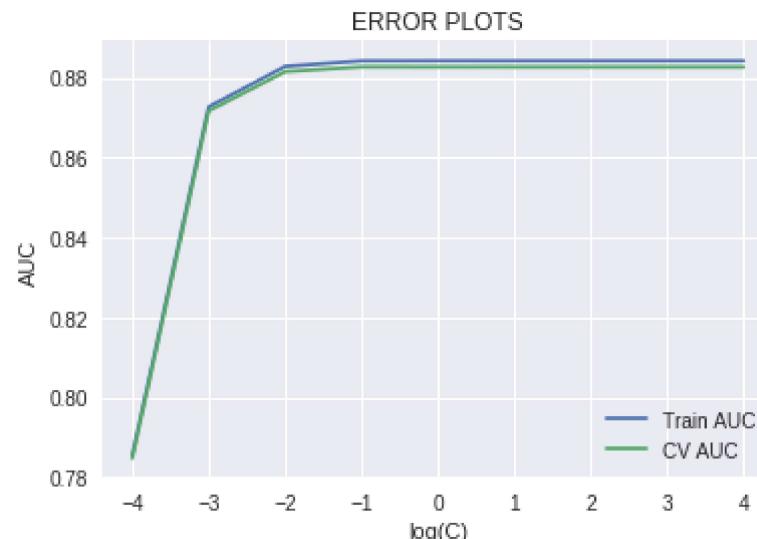
```
lrModel=LogisticRegression(penalty='l2')
hyper_parameters={'C':[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]}
odel=GridSearchCV(lrModel,hyper_parameters,scoring='roc_auc',cv=3)
odel.fit(X_train_tfidf_sent_vectors,y_train)

_C_values=[10**-4,10**-3,10**-2,10**-1,10**0,10**1,10**2,10**3,10**4]
train_auc= model.cv_results_['mean_train_score']
cv_auc = model.cv_results_['mean_test_score']

=np.log10(C_values)

lt.plot(C, train_auc, label='Train AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='darkblue')

lt.plot(C, cv_auc, label='CV AUC')
this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(alpha_values,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='darkorange')
lt.legend()
lt.xlabel("log(C)")
lt.ylabel("AUC")
lt.title("ERROR PLOTS")
lt.show()
```



```
In [0]: odel.best_estimator_
```

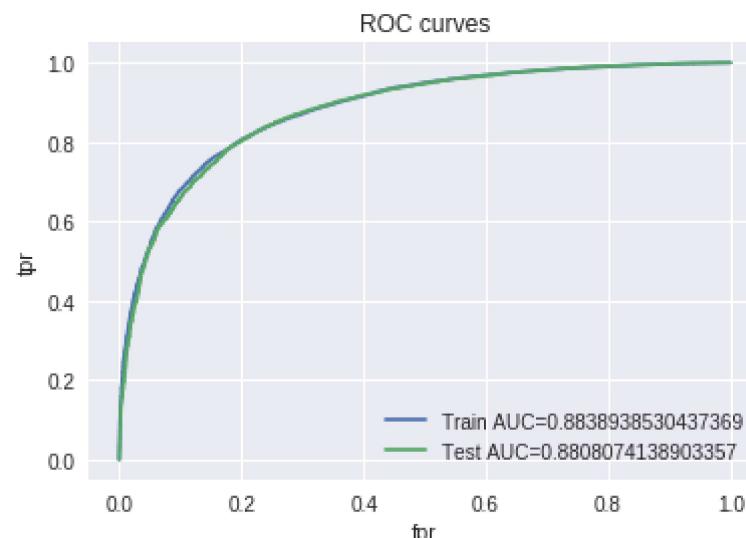
```
Out[0]: LogisticRegression(C=0.1, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='warn',
                           n_jobs=None, penalty='l2', random_state=None, solver='warn',
                           tol=0.0001, verbose=0, warm_start=False)
```

```
In [0]: est_C=0.1
```

```
In [0]: rModel=LogisticRegression(penalty='l2',C=best_C)
rModel.fit(X_train_tfidf_sent_vectors,y_train)

train_fpr,train_tpr,thresholds=roc_curve(y_train,lrModel.predict_proba(X_train_tfidf_sent_vectors)[:,1])
est_fpr,test_tpr,thresholds=roc_curve(y_test,lrModel.predict_proba(X_test_tfidf_sent_vectors)[:,1])

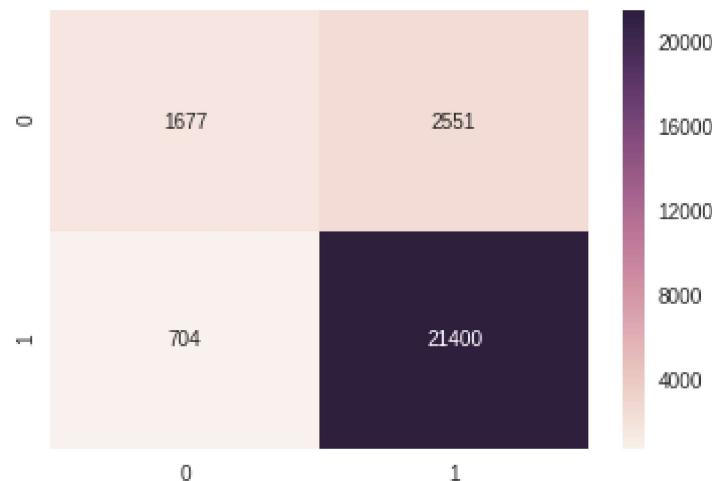
lt.plot(train_fpr,train_tpr,label='Train AUC='+str(auc(train_fpr,train_tpr)))
lt.plot(test_fpr,test_tpr,label='Test AUC='+str(auc(test_fpr,test_tpr)))
lt.legend()
lt.xlabel('fpr')
lt.ylabel('tpr')
lt.title('ROC curves')
lt.show()
```



In [0]: confusion matrix using seaborn heatmap

```
_pred=lrModel.predict(X_test_tfidf_sent_vectors)
m=confusion_matrix(y_test,y_pred)
ns.heatmap(cm,annot=True,fmt='.5g')
```

Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff8ec2cbe48>



[6] Conclusions

In [125]: <http://zetcode.com/python/prettytable/>

```
from prettytable import PrettyTable

= PrettyTable()

.field_names = ["Vectorizer", "Regularization", "Best C", "Test AUC"]

.add_row(["BOW", "L1", 1, 0.93])
.add_row(["BOW", "L2", 0.1, 0.94])
.add_row(["TFIDF", "L1", 1, 0.95])
.add_row(["TFIDF", "L2", 10, 0.96])
.add_row(["AvgW2V", "L1", 1, 0.9])
.add_row(["AvgW2V", "L2", 0.1, 0.9])
.add_row(["Tfidf W2V", "L1", 1, 0.88])
.add_row(["Tfidf W2V", "L2", 0.1, 0.88])

print(x)
```

Vectorizer	Regularization	Best C	Test AUC
BOW	L1	1	0.93
BOW	L2	0.1	0.94
TFIDF	L1	1	0.95
TFIDF	L2	10	0.96
AvgW2V	L1	1	0.9
AvgW2V	L2	0.1	0.9
Tfidf W2V	L1	1	0.88
Tfidf W2V	L2	0.1	0.88