

sklearn SGDRegressor

```
In [0]: import warnings
warnings.filterwarnings("ignore")
from sklearn.datasets import load_boston
from random import seed
from random import randrange
from csv import reader
from math import sqrt
from sklearn import preprocessing
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from prettytable import PrettyTable
from sklearn.linear_model import SGDRegressor
from sklearn import preprocessing
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
```

```
In [0]: #Loading the boston dataset from sklearn
X = load_boston().data
Y = load_boston().target
```

```
In [0]: #splitting the dataset into train and test dataset
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.3,random_state=0)
```

```
In [0]: #standardizing the data
scaler=preprocessing.StandardScaler()
scaler.fit(X_train)
X_train=scaler.transform(X_train)
X_test=scaler.transform(X_test)
```

```
In [0]: #sklearn SGDRegressor model without regularization  
clf = SGDRegressor(penalty=None,alpha=0,max_iter=500)  
#clf = SGDRegressor()  
clf.fit(X_train, Y_train)
```

```
Out[0]: SGDRegressor(alpha=0, average=False, early_stopping=False, epsilon=0.1,  
eta0=0.01, fit_intercept=True, l1_ratio=0.15,  
learning_rate='invscaling', loss='squared_loss', max_iter=500,  
n_iter=None, n_iter_no_change=5, penalty=None, power_t=0.25,  
random_state=None, shuffle=True, tol=None, validation_fraction=0.1,  
verbose=0, warm_start=False)
```

```
In [0]: #weight vector obtained by sklearn SGDRegressor model  
sklearn_W=clf.coef_  
sklearn_W
```

```
Out[0]: array([-1.01498388,  1.04914638,  0.08204709,  0.63096429, -1.87417512,  
2.69913176, -0.27455695, -3.10496178,  2.10098962, -1.88272296,  
-2.26244145,  0.58231771, -3.44130655])
```

```
In [0]: #intercept for the model obtained by sklearn SGDRegressor  
sklearn_b=clf.intercept_  
sklearn_b
```


```
Out[0]: array([22.74569767])
```

```
In [0]: #mean squared error  
Y_pred=clf.predict(X_test)  
sklearn_MSE=mean_squared_error(Y_test,Y_pred)
```

```
In [0]: sklearn_MSE
```

```
Out[0]: 27.173522895888894
```

```
In [0]: #plotting predicted values VS actual values  
plt.plot(Y_pred,Y_test,linestyle='',marker='o')  
plt.xlabel('predicted')  
plt.ylabel('actual')  
plt.show()
```



My implementation of SGDRegressor

```
In [0]: #Loading the dataset
        boston=load_boston()
        X=boston.data
        Y=boston.target

        #splitting the dataset into train,test datasets
        X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.3,random_state=0)

        #feature scaling
        scaler=preprocessing.StandardScaler().fit(X_train)
        X_train=scaler.transform(X_train)
        X_test=scaler.transform(X_test)
```

```
In [0]: print('Shape of X_train',X_train.shape,'\nShape of X_test',X_test.shape)
```

```
Shape of X_train (354, 13)
Shape of X_test (152, 13)
```

```

In [0]: iterations=500 #max no of iterations
        sample_size=150 #sample size of data to consider during each iteration
        W=np.random.normal(0.0,1.0,size=13) #initializing the weight vector with random values from normal distribution(mean=0
        and std-dev=1)
        b=np.random.normal(0.0,1.0,size=1) #initializing intercept term with a random value from normal distribution
        mse=[] #empty list for storing mse in each iteration
        lr=0.01 #learning rate

        for itr in range(iterations):
            #generating random numbers to be used as index for sampling
            idx=np.random.choice(np.arange(len(X_train)),size=sample_size,replace=False)
            X_sample=X_train[idx]
            Y_sample=Y_train[idx]
            #predicted values
            Y_pred=np.dot(X_sample,W)+b
            mse.append(mean_squared_error(Y_sample,Y_pred))
            if(itr!=0):
                if(abs(mse[itr]-mse[itr-1])>=0.1):
                    for i in range(len(X_sample)):
                        yhat=np.dot(W.T,X_sample[i])+b
                        W=W-lr*(-2)*X_sample[i]*(Y_sample[i]-yhat)
                        b=b-lr*(-2)*(Y_sample[i]-yhat)
                    lr=lr/2
                else:
                    break

```

```

In [0]: #MSE for my implementation
        Y_pred=np.dot(X_test,W)+b
        my_MSE=mean_squared_error(Y_test,Y_pred)
        my_MSE

```

Out[0]: 27.09624954954564

```

In [0]: #weights obtained by my implementation
        my_W=W
        print(my_W)

```

```


[-1.02467402  0.80761971 -0.34990453  0.43953371 -1.2261729  2.63689706
 0.25720897 -2.33834904  1.50005187 -1.04157106 -2.0133921  0.72474908
-3.99110089]

```

```
In [0]: #intercept obtained by my implementation
my_b=b
print(b)

[22.72308662]
```

```
In [0]: #predicted values vs actual values
plt.plot(Y_pred,Y_test,linestyle='',marker='o')
plt.xlabel('predicted')
plt.ylabel('actual')
plt.show()
```



Comparing my implementation and sklearn implementation of SGDRegressor

```
In [0]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Weights", "sklearn implementation", "my implementation"]

for i in range(13):
    x.add_row(["W"+str(i+1),sklearn_W[i],my_W[i]])

x.add_row(["b",sklearn_b,my_b])
print(x)
```

Weights	sklearn implementation	my implementation
W1	-1.0149838764122392	-1.024674018911933
W2	1.0491463807290133	0.8076197058072463
W3	0.08204709287137066	-0.34990453180233544
W4	0.6309642912826718	0.4395337060854564
W5	-1.8741751216435873	-1.226172897391543
W6	2.6991317567594706	2.636897064116853
W7	-0.2745569508530523	0.25720896610251
W8	-3.104961779342055	-2.3383490370415436
W9	2.10098961570462	1.5000518747058718
W10	-1.8827229635000406	-1.041571061518702
W11	-2.2624414505209454	-2.013392096886359
W12	0.5823177105483278	0.7247490768930089
W13	-3.441306549423374	-3.991100894226346
b	[22.74569767]	[22.72308662]

Comparing MSE for my implementation and sklearn implementation

```
In [0]: print('sklearn implementation MSE: ',sklearn_MSE)
        print('my implementation MSE: ',my_MSE)
```

```
sklearn implementation MSE: 27.173522895888894
my implementation MSE: 27.09624954954564
```

RMSE for my implementation and sklearn implementation

```
In [0]: sklearn_RMSE=np.sqrt(sklearn_MSE)
        my_RMSE=np.sqrt(my_MSE)
```

```
In [0]: print('sklearn implementation RMSE: ',sklearn_RMSE)
        print('my implementation RMSE: ',my_RMSE)
```

```
sklearn implementation RMSE: 5.212822929650391
my implementation RMSE: 5.205405800660083
```